

Clase 8

Consigna: Por cada ejercicio, escribir el código y agregar una captura de pantalla del resultado obtenido.

Diccionario de datos:

<https://www.kaggle.com/datasets/rohanrao/formula-1-world-championship-1950-2020?select=results.csv>

1. Crear la siguientes tablas externas en la base de datos f1 en hive:

- a. driver_results (driver_forename, driver_surname, driver_nationality, points)
- b. constructor_results (constructorRef, cons_name, cons_nationality, url, points)

```
create database f1;

create external table driver_results(
    driver_forename STRING,
    driver_surname STRING,
    driver_nationality STRING,
    points INT
)
-- row format delimited
-- fields terminated by ','
-- location '/tables/external/tripdata';

create external table constructor_results(
    constructorRef STRING,
    cons_name STRING,
    cons_nationality STRING,
    url STRING,
    points INT
)
-- row format delimited
-- fields terminated by ','
-- location '/tables/external/tripdata';
;
```

2. En Hive, mostrar el esquema de driver_results y constructor_results

```
describe formatted driver_results;  
describe formatted constructor_results;
```

```
hive> create external table constructor_results(  
  >     constructorRef STRING,  
  >     cons_name STRING,  
  >     cons_nationality STRING,  
  >     url STRING,  
  >     points INT  
  > )  
  > -- row format delimited  
  > -- fields terminated by ','  
  > -- location '/tables/external/tripdata';  
  > ;
```

OK

Time taken: 0.326 seconds

```
hive> describe formatted constructor_results;
```

OK

#	col_name	data_type	comment
	constructorref	string	
	cons_name	string	
	cons_nationality	string	
	url	string	
	points	int	

```
hive> describe formatted driver_results;
```

OK

#	col_name	data_type	comment
	driver_forename	string	
	driver_surname	string	
	driver_nationality	string	
	points	int	

Detailed Table Information

Database:	f1
Owner:	hadoop
CreateTime:	Thu May 23 22:30:44 ART 2024
LastAccessTime:	UNKNOWN
Retention:	0
Location:	hdfs://172.17.0.2:9000/user/hive/warehouse/f1.db/driver_results
Table Type:	EXTERNAL_TABLE
Table Parameters:	

3. Crear un archivo .bash que permita descargar los archivos mencionados abajo e ingestarlos en HDFS:

results.csv

<https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/results.csv>

drivers.csv

<https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/drivers.csv>

constructors.csv

<https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/constructors.csv>

races.csv

<https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/races.csv>

```
#!/bin/bash

# Clean landing directory and get data
rm /home/hadoop/landing/*

wget -P /home/hadoop/landing "https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/results.csv"
wget -P /home/hadoop/landing "https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/drivers.csv"
wget -P /home/hadoop/landing "https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/constructors.csv"
wget -P /home/hadoop/landing "https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/races.csv"

# Clean HDFS directory and put the data from landing into Hadoop
/home/hadoop/hadoop/bin/hdfs dfs -rm /ingest/*
/home/hadoop/hadoop/bin/hdfs dfs -put /home/hadoop/landing/* /ingest
```

4. Generar un archivo .py que permita, mediante Spark:

- a. insertar en la tabla driver_results los corredores con mayor cantidad de puntos en la historia.
- b. insertar en la tabla constructor_result quienes obtuvieron más puntos en el Spanish Grand Prix en el año 1991

```
from pyspark.context import SparkContext
from pyspark.sql.session import SparkSession
from pyspark.sql import HiveContext

# Configuramos...
sc = SparkContext('local')
spark = SparkSession(sc)
hc = HiveContext(sc)

# Traemos tablas
df1 = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/results.csv")
df2 = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/drivers.csv")
df3 = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/constructors.csv")
df4 = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/races.csv")

# Preparamos la primera transformacion
df1.createOrReplaceTempView("results")
df2.createOrReplaceTempView("drivers")

driver_results = spark.sql("""
    SELECT
        CAST(d.forename AS STRING) AS driver_forename,
        CAST(d.surname AS STRING) AS driver_surname,
        CAST(d.nationality AS STRING) AS driver_nationality,
        CAST(r.points AS INT) AS points
    FROM drivers d
    INNER JOIN results r ON d.driverId = r.driverId
    WHERE r.points != 0
""")

# Preparamos la segunda transformacion
df3.createOrReplaceTempView("constructors")
df4.createOrReplaceTempView("races")

SpanishGP = spark.sql("""
    SELECT
        CAST(c.constructorRef AS STRING) AS constructorref,
        CAST(c.name AS STRING) AS cons_name,
        CAST(c.nationality AS STRING) AS cons_nationality,
        CAST(c.url AS STRING) AS url,
        CAST(r.points AS INT) AS points
    FROM constructors c
    INNER JOIN results r ON c.constructorId = r.constructorId
    INNER JOIN races ra ON ra.raceId = r.raceId
    WHERE ra.circuitID IN (4, 12, 26, 45, 49, 67) AND r.points != 0 AND ra.year = 1991
""")

SpanishGP.show(3)

# Insertamos las tablas
driver_results.write.insertInto("f1.driver_results")
SpanishGP.write.insertInto("f1.constructor_results")
```

5. Realizar un proceso automático en Airflow que orqueste los archivos creados en los puntos 3 y 4. Correrlo y mostrar una captura de pantalla (del DAG y del resultado en la base de datos)

```
from pyspark.context import SparkContext
from pyspark.sql.session import SparkSession
from pyspark.sql import HiveContext

# Configuramos...
sc = SparkContext('local')
spark = SparkSession(sc)
hc = HiveContext(sc)

# Traemos tablas
df1 = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/results.csv")
df2 = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/drivers.csv")
df3 = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/constructors.csv")
df4 = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/races.csv")

# Preparamos la primera transformacion
df1.createOrReplaceTempView("results")
df2.createOrReplaceTempView("drivers")

driver_results = spark.sql("""
    SELECT
        CAST(d.forename AS STRING) AS driver_forename,
        CAST(d.surname AS STRING) AS driver_surname,
        CAST(d.nationality AS STRING) AS driver_nationality,
        CAST(r.points AS INT) AS points
    FROM drivers d
    INNER JOIN results r ON d.driverId = r.driverId
    WHERE r.points != 0
""")

# Preparamos la segunda transformacion
df3.createOrReplaceTempView("constructors")
df4.createOrReplaceTempView("races")

SpanishGP = spark.sql("""
    SELECT
        CAST(c.constructorRef AS STRING) AS constructorref,
        CAST(c.name AS STRING) AS cons_name,
        CAST(c.nationality AS STRING) AS cons_nationality,
        CAST(c.url AS STRING) AS url,
        CAST(r.points AS INT) AS points
    FROM constructors c
    INNER JOIN results r ON c.constructorId = r.constructorId
    INNER JOIN races ra ON ra.raceId = r.raceId
    WHERE ra.circuitID IN (4, 12, 26, 45, 49, 67) AND r.points != 0 AND ra.year = 1991
""")

SpanishGP.show(3)

# Insertamos las tablas
driver_results.write.insertInto("f1.driver_results")
SpanishGP.write.insertInto("f1.constructor_results")
```

